

# Anytime Incremental $\rho$ POMDP Planning in Continuous Spaces

Presented by: Ron Benchetrit  
Advisor: Vadim Indelman

Technion – Israel Institute of Technology

28 April 2025



# Outline

- 1 Introduction
- 2 Background
- 3 Solving POMDPs
- 4  $\rho$ POMCPOW - Challenges and Solutions
- 5 Conclusions

# Autonomous Agents

- Autonomous agents appear in a wide range of domains:




# Autonomous Agents

- Autonomous agents must:
  - **Perceive** the world through noisy, partial observations
  - **Act** under uncertainty in outcomes and dynamics
  - **Plan** to achieve long-term objectives
- This talk focuses on **sequential decision-making under uncertainty**.
- Specifically, planning in the **Partially Observable Markov Decision Process (POMDP)** and its extension, the  $\rho$ POMDP.

# MDP – Example

- Imagine a Roomba-like robot operating in a room.
- Its goal is to navigate to a designated target area.
- The robot acts in the environment by actuating its motors.
- However, due to motor errors, slippage, etc., the outcome of an action may not be as expected.
- We want the robot to reach the goal quickly and safely, avoiding obstacles along the way.



# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.

# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.

# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.
- **Definition:**  $\mathcal{A}$  is the action space.



# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.
- **Definition:**  $\mathcal{A}$  is the action space.
- *Example:* move forward, turn left, turn right.

# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.
- **Definition:**  $\mathcal{A}$  is the action space.
- *Example:* move forward, turn left, turn right.
- **Definition:**  $\mathcal{T}(s' | a, s)$  is the transition model — the probability of reaching  $s'$  after taking action  $a$  in state  $s$ .

# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.
- **Definition:**  $\mathcal{A}$  is the action space.
- *Example:* move forward, turn left, turn right.
- **Definition:**  $\mathcal{T}(s' | a, s)$  is the transition model — the probability of reaching  $s'$  after taking action  $a$  in state  $s$ .
- *Example:* A linear-Gaussian motion model that accounts for motor noise and slippage.

# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.
- **Definition:**  $\mathcal{A}$  is the action space.
- *Example:* move forward, turn left, turn right.
- **Definition:**  $\mathcal{T}(s' | a, s)$  is the transition model — the probability of reaching  $s'$  after taking action  $a$  in state  $s$ .
- *Example:* A linear-Gaussian motion model that accounts for motor noise and slippage.
- **Definition:**  $\mathcal{R}(s, a)$  is the reward function.

# MDP – Definition

A Markov Decision Process (MDP) is defined by the tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R})$ :

- **Definition:**  $\mathcal{S}$  is the state space.
- *Example:* All possible robot positions and headings in the room.
- **Definition:**  $\mathcal{A}$  is the action space.
- *Example:* move forward, turn left, turn right.
- **Definition:**  $\mathcal{T}(s' | a, s)$  is the transition model — the probability of reaching  $s'$  after taking action  $a$  in state  $s$ .
- *Example:* A linear-Gaussian motion model that accounts for motor noise and slippage.
- **Definition:**  $\mathcal{R}(s, a)$  is the reward function.
- *Example:* +100 for reaching the goal, −100 for hitting an obstacle, and −1 for each time step.

# MDP – Objective

The agent aims to choose actions that **maximize expected cumulative reward**:

$$\mathbb{E} \left[ \sum_{t=0}^{\infty} r(s_t, a_t) \right]$$

A solution to an MDP is a **policy**  $\pi : \mathcal{S} \rightarrow \mathcal{A}$  that maps states to actions.

- **Value function:**

$$V^{\pi}(s) = \mathbb{E} \left[ \sum_{t=0}^{\infty} r(s_t, a_t) \mid s_0 = s, a_t = \pi(s_t) \right]$$

- **Action-value function:**

$$Q^{\pi}(s, a) = \mathbb{E}_{s' \sim T(\cdot | s, a)} [r(s, a) + V^{\pi}(s')] , \quad V^{\pi}(s) = Q^{\pi}(s, \pi(s))$$

# POMDP – Example

- In the MDP setting, the robot always knew its exact position.
- Now, its position is **unknown** — it only starts with a rough guess.
- The robot uses a **laser rangefinder** to measure distances to walls.
- These measurements are **noisy and ambiguous** — multiple locations can produce similar readings.
- To navigate effectively, the robot must estimate its position using both its actions and sensor data.

# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.



# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.
- **Definition:**  $\mathcal{O}$  is the observation space.

# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.
- **Definition:**  $\mathcal{O}$  is the observation space.
- *Example:* All possible laser rangefinder readings the robot can receive.

# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.
- **Definition:**  $\mathcal{O}$  is the observation space.
- *Example:* All possible laser rangefinder readings the robot can receive.
- **Definition:**  $\mathcal{Z}(o \mid s)$  is the observation model — the probability of observing  $o$  in state  $s$ .

# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.
- **Definition:**  $\mathcal{O}$  is the observation space.
- *Example:* All possible laser rangefinder readings the robot can receive.
- **Definition:**  $\mathcal{Z}(o \mid s)$  is the observation model — the probability of observing  $o$  in state  $s$ .
- *Example:* Noisy distance measurements that may be consistent with several positions.

# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.
- **Definition:**  $\mathcal{O}$  is the observation space.
- *Example:* All possible laser rangefinder readings the robot can receive.
- **Definition:**  $\mathcal{Z}(o \mid s)$  is the observation model — the probability of observing  $o$  in state  $s$ .
- *Example:* Noisy distance measurements that may be consistent with several positions.
- **Definition:**  $b_0$  is the initial belief - a distribution over states.

# POMDP – Definition

A Partially Observable Markov Decision Process (POMDP) extends an MDP to handle partial observability. It is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ , and  $\mathcal{R}$  are the same as in MDP.
- **Definition:**  $\mathcal{O}$  is the observation space.
- *Example:* All possible laser rangefinder readings the robot can receive.
- **Definition:**  $\mathcal{Z}(o \mid s)$  is the observation model — the probability of observing  $o$  in state  $s$ .
- *Example:* Noisy distance measurements that may be consistent with several positions.
- **Definition:**  $b_0$  is the initial belief - a distribution over states.
- *Example:* The robot's initial guess about its position, e.g., a uniform distribution over the room.

# Belief – Definition

- In MDPs, the agent chooses actions based on the fully observable state.
- In POMDPs, the state is hidden, so the agent must act based on what it knows so far.
- To handle this, the agent maintains a **belief**  $b_t$ : a probability distribution over possible states at time  $t$ .
- The belief is updated over time using the agent's history of actions and observations:

$$h_t = (b_0, a_0, o_1, \dots, a_{t-1}, o_t)$$

- The belief at time  $t$  is defined as:

$$b_t(s) = \mathbb{P}(s \mid h_t)$$

# POMDP – Objective

As in MDPs, the agent aims to **maximize expected cumulative reward**:

$$\mathbb{E} \left[ \sum_{t=0}^{\infty} r(s_t, a_t) \right]$$

Since the state is hidden, actions are chosen based on the agent's **belief**. A solution is a **policy**  $\pi : \mathcal{B} \rightarrow \mathcal{A}$  mapping beliefs to actions.

- **Value function:**

$$V^{\pi}(b) = \mathbb{E} \left[ \sum_{t=0}^{\infty} r(s_t, a_t) \mid b_0 = b, a_t = \pi(b_t) \right]$$

- **Action-value function:**

$$Q^{\pi}(b, a) = \mathbb{E}_{s \sim b, s' \sim T(\cdot|s,a), o \sim O(\cdot|s')} [r(s, a) + V^{\pi}(b')]$$

where  $b'$  is the updated belief after taking  $a$  and observing  $o$ .



The belief is updated recursively using the Bayes filter:

$$b_t(s_t) = \eta \cdot \mathcal{Z}(o_t \mid s_t) \int_S \mathcal{T}(s_t \mid s_{t-1}, a_{t-1}) b_{t-1}(s_{t-1}) ds_{t-1}$$

where  $\eta$  is a normalization constant.


Special cases of the Bayes filter include:

- **Kalman Filter** — for linear-Gaussian systems.
- **Particle Filter** — for non-linear, non-Gaussian systems.

# Particle Filter


- A particle filter approximates the belief using a **set of weighted samples** (particles).
- Each particle represents a possible state and has an associated weight.

$$b(s) \approx \sum_{i=1}^N w_i \cdot \delta(s - s_i)$$



- This representation is highly flexible and well-suited for capturing complex, multimodal beliefs.
- As the number of particles  $N$  **increases**, the approximation becomes more **accurate**.

# Particle Filter – Update



- Now, suppose the robot's task is to **localize itself** in the room.
- There is no designated goal location — the objective is to **reduce uncertainty**.
- Examples: active localization, informative path planning, active learning, and active SLAM.
- Such tasks cannot be naturally expressed using the standard  $R(s, a)$ .
- Instead, we define a **belief-dependent reward function**  $\rho(b, a)$  — for example, negative entropy.

# $\rho$ POMDP – Definition

A  $\rho$ POMDP is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0, \rho)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ ,  $\mathcal{O}$ ,  $\mathcal{Z}$ , and  $b_0$  are defined as in a standard POMDP.

# $\rho$ POMDP – Definition

A  $\rho$ POMDP is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0, \rho)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ ,  $\mathcal{O}$ ,  $\mathcal{Z}$ , and  $b_0$  are defined as in a standard POMDP.
- **Definition:**  $\rho$  is a belief dependent reward, i.e.,  $\rho(b, a)$ .

# $\rho$ POMDP – Definition

A  $\rho$ POMDP is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0, \rho)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ ,  $\mathcal{O}$ ,  $\mathcal{Z}$ , and  $b_0$  are defined as in a standard POMDP.
- **Definition:**  $\rho$  is a belief dependent reward, i.e.,  $\rho(b, a)$ .
- *Example:* Information-theoretic rewards such as negative entropy or information gain.

# $\rho$ POMDP – Definition

A  $\rho$ POMDP is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0, \rho)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ ,  $\mathcal{O}$ ,  $\mathcal{Z}$ , and  $b_0$  are defined as in a standard POMDP.
- **Definition:**  $\rho$  is a belief dependent reward, i.e.,  $\rho(b, a)$ .
- *Example:* Information-theoretic rewards such as negative entropy or information gain.
- This formulation naturally captures tasks centered on **active information gathering**.



A  $\rho$ POMDP is defined by the tuple:

$$(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O}, \mathcal{Z}, b_0, \rho)$$

- $\mathcal{S}$ ,  $\mathcal{A}$ ,  $\mathcal{T}$ ,  $\mathcal{O}$ ,  $\mathcal{Z}$ , and  $b_0$  are defined as in a standard POMDP.
- **Definition:**  $\rho$  is a belief dependent reward, i.e.,  $\rho(b, a)$ .
- *Example:* Information-theoretic rewards such as negative entropy or information gain.
- This formulation naturally captures tasks centered on **active information gathering**.
- In practice,  $\rho$  can combine both belief-dependent and state-dependent rewards:

$$\rho(b, a) = \mathbb{E}_{s \sim b}[r(s, a)] + \mathcal{I}(b)$$

In  $\rho$ POMDPs, the agent aims to **maximize expected cumulative belief-dependent reward**:

$$\mathbb{E} \left[ \sum_{t=0}^{\infty} \rho(b_t, a_t) \right]$$

Here, rewards depend directly on the **belief** and **action**. A solution is a **policy**  $\pi : \mathcal{B} \rightarrow \mathcal{A}$  mapping beliefs to actions.

# Solving POMDPs – Computation Complexity

Solving POMDPs is a challenging task due to the following reasons:


- Curse of dimensionality
- Curse of history
- Continuous state space
- Continuous action space
- Continuous observation space

# Offline vs. Online Planning

- **Offline methods** aim to compute a policy  $\pi : \mathcal{B} \rightarrow \mathcal{A}$  in advance.
  - Require solving for the entire belief space.
  - Computationally intractable for large or continuous domains.
  - Examples: value iteration, point-based solvers (PBVI).
- **Online methods** compute the next action at each time step, based on the current belief.
  - Avoid planning for unreachable beliefs.
  - More scalable in high-dimensional or continuous settings.
  - Examples: POMCP [6], DESPOT [7], POMCPOW [8].


# Online Tree Search

- Online methods plan by building a search tree rooted at the current belief — the **belief tree**.
- Nodes represent beliefs; edges correspond to actions and observations.
- The tree is expanded by simulating action-observation trajectories.
- The best action is selected based on estimated values at the root's children.
- The chosen action is executed in the real environment, an observation is received, the belief is updated, and the process repeats.




# Monte Carlo Tree Search

- Monte Carlo Tree Search (MCTS) is a widely used planning algorithm.
- It uses random trajectory simulations to estimate the value of actions from the current belief.
- MCTS builds the tree incrementally using four key steps:
- **Selection:** Traverse the tree to select a promising node.
- **Expansion:** Add one or more child nodes.
- **Simulation:** Simulate a rollout from the new node to estimate return.
- **Backpropagation:** Propagate the return up the tree to update value estimates.




# State Simulators vs. Belief Simulators

- SOTA online POMDP solvers use a **particle-based** approach to represent the belief.
- These solvers can be broadly categorized as **state simulators** or **belief simulators**.



# State Simulators


- Simulate state trajectories directly.
- Each time a belief node is visited, a new particle is added.
- Particles accumulated across visits.
- Representation improves at frequently visited (and promising) nodes.
- Belief-dependent rewards are more challenging to compute.
- Examples: POMCP [6], POMCPOW [8], DESPOT [7], LABECOP [5].





# Belief Simulators

- Treat belief states as explicit nodes in a belief-MDP.
- Fixed set of particles created once upon expansion.
- Representation is static, wasting computation resources.
- Easy belief-dependent rewards computation.
- Examples: PFT-DPW [8], AdaOPS [10].



Setting	Belief Simulator	State Simulator
Discrete	$\rho$ beliefUCT [9]	$\rho$ POMCP [9]
Continuous	PFT-DPW [8] IPFT [4] AI-FSSS [1] SITH-PFT [11]	

Setting	Belief Simulator	State Simulator
Discrete	$\rho$ beliefUCT [9]	$\rho$ POMCP [9]
Continuous	PFT-DPW [8] IPFT [4] AI-FSSS [1] SITH-PFT [11]	$\rho$ POMCPOW (ours)

- We aim to extend POMCPOW — a state-of-the-art state-simulator for POMDPs — to the  $\rho$ POMDP setting.
- This introduces several key challenges:
  - **Value estimation:** Running averages used in POMCPOW are incompatible with belief-dependent rewards.
  - **Belief representation:** Particle accumulation during visits does not ensure sufficient belief quality across the tree.
  - **Reward recomputation:** Belief updates require frequent and costly re-evaluation of the reward function.

In the next slides, we address each of these challenges in turn.

# Challenge: Value Estimation

- In POMCPOW, node values are estimated using a running average of sampled rewards:

$$\hat{Q}(h, a) = \frac{1}{N(ha)} \sum_{i=1}^{N(ha)} R_i$$

- This is well-suited for state-based rewards, where the expected reward depends only on the state.
- In  $\rho$ POMDPs, rewards are **non-linear functions of the belief**, which evolves as new particles are added.
- Averaging rewards across outdated beliefs leads to **biased and inconsistent** value estimates.

# Solution: Last-Value Update (LVU)

- To avoid mixing outdated rewards, we adopt the **Last-Value Update (LVU)** framework proposed by [9].
- In LVU, each node stores a value estimate based only on the most recent rewards from its children.

**Value:**

$$\hat{V}(h) = \frac{1}{N(h)} \left[ \text{Rollout}(h) + \sum_{a \in \text{Ch}(h)} N(ha) \hat{Q}(ha) \right]$$

**Action-value:**

$$\hat{Q}(ha) = \frac{1}{N(ha)} \sum_{o \in \text{Ch}(ha)} N(hao) \left[ \hat{r}(hao) + \gamma \hat{V}(hao) \right]$$

# Our Contribution: Incremental LVU

- Standard LVU updates require summing over all children:  $O(n)$  time.
- This is costly in large trees, especially in continuous domains.
- We propose an **incremental update** that maintains correctness with just  $O(1)$  computation.

## Incremental value update:

$$\hat{V}(h) \leftarrow \hat{V}(h) + \frac{1}{N(h)} \left[ N(ha') \hat{Q}(ha') - (N(ha') - 1) \hat{Q}^{\text{prev}}(ha') - \hat{V}(h) \right]$$

## Incremental action-value update:

$$\begin{aligned} \hat{Q}(ha) \leftarrow \hat{Q}(ha) + \frac{1}{N(ha)} & \left[ N(hao') (\hat{\rho}(hao') + \gamma \hat{V}(hao')) \right. \\ & \left. - (N(hao') - 1) (\hat{\rho}^{\text{prev}}(hao') + \gamma \hat{V}^{\text{prev}}(hao')) - \hat{Q}(ha) \right] \end{aligned}$$

# Challenge: Belief Representation

- Beliefs are typically represented using a particle-based approximation.
- In state simulators, particles are added incrementally during node visitation.
- This improves belief representation at frequently visited nodes.
- However, less-visited nodes may remain poorly represented.
- For belief-dependent rewards, we need good belief quality **throughout the tree**, not just locally.



# Solution: Anytime Belief Refinement

- We ensure that belief quality improves **over time and throughout the tree**.
- Our approach provides a theoretical guarantee: under consistent selection strategies, **every node** is visited increasingly often.
- This enables belief refinement even in deep or rarely explored parts of the tree.
- Further details and formal results are available in the paper.

# Challenge: Reward Re-computation

- In  $\rho$ POMDPs, rewards depend on the belief — not just the state.
- As the belief evolves (e.g., when new particles are added), the reward must be updated accordingly.
- However, belief-dependent rewards are often:
  - **Non-linear** (e.g., entropy-based), making updates non-trivial.
  - **Expensive** to recompute from scratch.
- Major computational bottleneck in continuous domains.
- Want to update rewards **incrementally**, without full recomputation.
- We will focus in information-theoretic rewards.

# Solution: Incremental Shannon Entropy

- Shannon entropy is a common information-theoretic reward used to quantify belief uncertainty:

$$\hat{H}(b) = - \sum_{i=1}^N w_i \log w_i$$

- Recomputing it after every belief update is linear in the number of particles:  $O(N)$ .
- We derive an equivalent form that enables fast incremental updates:

$$\hat{H}(b) = - \frac{1}{\sum w_i} \sum w_i \log w_i + \log \sum w_i$$

- This formulation allows updating entropy in  $O(1)$  time after adding a new particle — by reusing cached terms.

# Boers Entropy Estimator

- Shannon entropy is simple but not well-suited for particle-based beliefs — especially in continuous spaces.
- The Boers estimator [2] is specifically designed for such settings, capturing local density and the shape of the belief.
- It is defined as:

$$\hat{H}(b') = \log \sum_{i=1}^N \mathcal{Z}(o \mid s'_i) w'_i - \sum_{i=1}^N w'_i \log \mathcal{Z}(o \mid s'_i) - \sum_{i=1}^N w'_i \log \left[ \sum_{j=1}^N \mathcal{T}(s'_i \mid s_j, a) w_j \right]$$

- The final term — a nested sum over particles — is the computational bottleneck:  $\mathcal{O}(N^2)$ .
- This term can be updated incrementally and efficiently.

# Incremental Update for Boers Estimator

- Recall the final (expensive) term in the Boers estimator:

$$\sum_{i=1}^N w'_i \log \underbrace{\left[ \sum_{j=1}^N \mathcal{T}(s'_i \mid s_j, a) \cdot w_j \right]}_{c_i}$$

- Adding a new particle  $s_{N+1}$  affects all  $c_i$ . We incrementally update them as:

$$\tilde{c}_i = \frac{\sum_{j=1}^N w_j}{\sum_{j=1}^{N+1} w_j} \cdot c_i + \mathcal{T}(s_i \mid s_{N+1}, a) \cdot \tilde{w}_{N+1}$$

- This reuses cached values and only requires computing one new transition column.
- Total cost drops from  $\mathcal{O}(N^2)$  to  $\mathcal{O}(N)$ .

# Experimental Setup

- We implemented  $\rho$ POMCPOW in Julia using the POMDPs.jl framework [3].
- Source code: [github.com/ronbenc/rhoPOMCPOW](https://github.com/ronbenc/rhoPOMCPOW)
- Compared against state-of-the-art solvers on two benchmark domains:
  - **Continuous 2D Light-Dark**
  - **Active Localization**
- We also evaluate the benefit of **incremental reward computation**.

# Benchmark Domains

- Both domains are set in a 2D continuous environment with:
  - Uncertain initial belief over the agent's state
  - Beacons that provide noisy observations (more accurate when closer)
  - Stochastic dynamics with 8 movement directions (unit circle) and a stay action that terminates the episode
- **Light-Dark:**
  - Task: navigate to a goal region from an uncertain start.
  - Sparse rewards — large bonus on success, heavy penalty on failure.
  - *Information gain* is used as reward shaping to encourage localization.
- **Active Localization:**
  - Task: actively reduce uncertainty about position.
  - Reward: pure *information gain*.
  - Environment includes obstacles that penalize collisions.
- *Results are averaged over 1000 runs with different seeds. We report mean return  $\pm$  standard error.*

# Results – Continuous 2D Light-Dark

- **Task:** Navigate to a goal region using beacon-based observations.
- $\rho$ POMCPOW outperforms due to its effective belief representation.
- Despite the added cost, incorporating *information gain* boosts planning effectiveness in  $\rho$ POMCPOW.

Algorithm	0.1s	0.2s	1.0s
$\rho$ POMCPOW <sup>†</sup>	<b>22.3</b> $\pm$ 1.2	<b>25.9</b> $\pm$ 1.1	<b>26.2</b> $\pm$ 1.1
POMCPOW	17.2 $\pm$ 1.4	17.5 $\pm$ 1.4	18.5 $\pm$ 0.9
IPFT <sup>†</sup>	-2.3 $\pm$ 1.8	6.4 $\pm$ 1.7	17.2 $\pm$ 1.2
PFT-DPW <sup>†</sup>	5.3 $\pm$ 1.6	13.4 $\pm$ 1.4	20.5 $\pm$ 1.0

<sup>†</sup> Uses information gain as reward shaping.



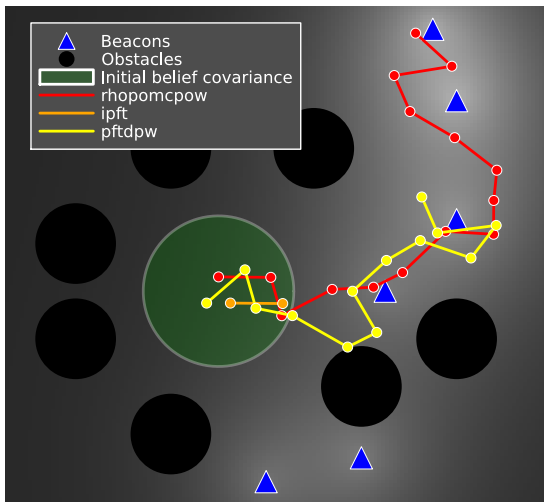
# Results – Active Localization

- **Task:** Actively localize using beacon observations while avoiding obstacles.
- $\rho$ POMCPOW **steadily improves** and outperforms as the planning time increases — thanks to its anytime refinement of beliefs.
- IPFT underperforms, likely due to its expensive particle reinvigoration mechanism.

Algorithm	0.1s	0.2s	1.0s
$\rho$ POMCPOW	29.0 $\pm$ 0.5	<b>38.1</b> $\pm$ 0.7	<b>45.9</b> $\pm$ 0.8
IPFT	27.1 $\pm$ 0.4	27.7 $\pm$ 0.4	27.0 $\pm$ 0.4
PFT-DPW	<b>36.7</b> $\pm$ 0.7	37.6 $\pm$ 0.7	38.7 $\pm$ 0.7


# Example Trajectory – Active Localization

- The agent actively reduces uncertainty about its position while avoiding obstacles.




# Effect of Incremental Reward Computation

- We compare  $\rho$ POMCPOW **with and without** incremental reward updates on the Light-Dark domain.
- Both variants use the same random seed to ensure identical search tree expansion.
- **Result:** Full recomputation becomes increasingly expensive as planning progresses.
- **Incremental updates significantly reduce planning time** — essential for  $\rho$ POMCPOW to be applicable in real-time scenarios.



- We introduced  **$\rho$ POMCPOW** — an anytime, online tree search algorithm for solving  $\rho$ POMDPs in continuous spaces.
- Our contributions include:
  - **Incremental LVU** for accurate and efficient value estimation.
  - **Anytime belief refinement** with theoretical guarantees on belief improvement.
  - **Incremental reward computation**, enabling scalable use of belief-dependent rewards.
- Empirical results demonstrate improved solution quality and planning efficiency over state-of-the-art methods.
- **Limitations:** convergence remains an open question; belief-dependent reward computation is still costly.

## THANK YOU!



## QUESTIONS?



M. Barenboim and V. Indelman.

Adaptive information belief space planning.

*In the 31st International Joint Conference on Artificial Intelligence and the 25th European Conference on Artificial Intelligence (IJCAI-ECAI), July 2022.*



Y. Boers, H. Driessen, A. Bagchi, and P. Mandal.

Particle filter based entropy.

*In 2010 13th International Conference on Information Fusion, pages 1–8, 2010.*



Maxim Egorov, Zachary N Sunberg, Edward Balaban, Tim A Wheeler, Jayesh K Gupta, and Mykel J Kochenderfer.

Pomdps. jl: A framework for sequential decision making under uncertainty.

*The Journal of Machine Learning Research, 18(1):831–835, 2017.*



Johannes Fischer and Omer Sahin Tas.

Information particle filter tree: An online algorithm for pomdps with belief-based rewards on continuous domains.

In *Intl. Conf. on Machine Learning (ICML)*, Vienna, Austria, 2020.



Marcus Hoerger and Hanna Kurniawati.

An on-line pomdp solver for continuous observation spaces.

In *IEEE Intl. Conf. on Robotics and Automation (ICRA)*, pages 7643–7649. IEEE, 2021.



David Silver and Joel Veness.

Monte-carlo planning in large pomdps.

In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 2164–2172, 2010.



Adhiraj Somani, Nan Ye, David Hsu, and Wee Sun Lee.

Despot: Online pomdp planning with regularization.

In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 13, pages 1772–1780, 2013.



Zachary Sunberg and Mykel Kochenderfer.

Online algorithms for pomdps with continuous state, action, and observation spaces.

In *Proceedings of the International Conference on Automated Planning and Scheduling*, volume 28, 2018.



Vincent Thomas, Jeremy Hutin, and Olivier Buffet.

Monte carlo information-oriented planning.

*arXiv preprint arXiv:2103.11345*, 2021.



Chenyang Wu, Guoyu Yang, Zongzhang Zhang, Yang Yu, Dong Li, Wulong Liu, and Jianye Hao.

Adaptive online packing-guided search for pomdps.

*Advances in Neural Information Processing Systems*, 34:28419–28430, 2021.



Andrey Zhitnikov, Ori Sztyglic, and Vadim Indelman.

No compromise in solution quality: Speeding up belief-dependent continuous pomdps via adaptive multilevel simplification.

*Intl. J. of Robotics Research*, 2024.